

U.S. NON-PROVISIONAL PATENT APPLICATION

THREE-DIMENSIONAL OBJECT RECONSTRUCTION FROM A VIDEO

CLAIM OF PRIORITY

[0001] This application is a continuation of U.S. Patent Application No. 17/734,244 titled “Three-Dimensional Object Reconstruction from a Video,” filed May 2, 2022 which is a continuation of U.S. Patent Application No. 16/945,455 titled “Three-Dimensional Object Reconstruction from a Video,” filed July 31, 2020, the entire contents of both applications are incorporated herein by reference.

TECHNICAL FIELD

[0002] The present invention relates to three-dimensional (3D) object reconstruction, and in particular, to a technique for constructing a 3D object from a video.

BACKGROUND

[0003] When we humans try to understand an image of an object, such as a duck, we instantly recognize a “duck”. We also instantly perceive and imagine the shape of the duck in the 3D world the duck’s appearance from other viewpoints. Furthermore, when we see the duck in a video, its 3D structure and deformation become even more apparent to us. Our ability to perceive the 3D structure of objects contributes vitally to our rich understanding of them.

[0004] While 3D perception is easy for humans, 3D reconstruction of deformable objects remains a very challenging problem in computer vision, especially for objects in the wild. For

learning-based algorithms, a bottleneck is the lack of supervision available for training. It is challenging to collect 3D annotations, such as 3D shape and camera pose, without limiting the domain (e.g., rigid objects, human bodies, and faces) for which 3D annotations can be captured in constrained environments. However, conventional approaches in the limited domains do not generalize well to non-rigid objects captured in naturalistic environments (e.g., animals). Due to constrained environments and limited annotations, it is very difficult to generalize the conventional approaches to the 3D construction of non-rigid objects (e.g., animals) from images and videos captured in the wild. There is a need for addressing these issues and/or other issues associated with the prior art.

SUMMARY

[0005] A 3D object reconstruction neural network system learns to predict a 3D representation of an object from a video that includes the object. An object in a video maintains temporal consistency, having a consistent shape and consistent texture across multiple frames. The temporal consistency of the object is exploited to reconstruct a dynamic 3D representation of the object from an unlabeled video. Texture, identity shape, and part correspondence invariance constraints may be applied to fine-tune the neural network system. The reconstruction technique generalizes well, particularly for non-rigid objects and the neural network system can inference in real time.

[0006] A method, computer readable medium, and system are disclosed for constructing a 3D representation of an object from a video. In an embodiment, a neural network model receives a video including images of the object captured from a camera pose and predicts a 3D shape representation of the object for a first image of the images based on a set of learned shape bases. The neural network model also predicts a texture flow for the first image and maps pixels from the first image to a texture space according to the texture flow to produce a texture image, where transfer of the texture image onto the 3D shape representation constructs a 3D object corresponding to the object in the first image.

BRIEF DESCRIPTION OF THE DRAWINGS

[0007] The present systems and methods for three-dimensional (3D) object reconstruction from a video are described in detail below with reference to the attached drawing figures, wherein:

[0008] Figure **1A** illustrates a block diagram of an example 3D object reconstruction system suitable for use in implementing some embodiments of the present disclosure.

[0009] Figure **1B** illustrates flowchart of a method for reconstructing a 3D representation of an object using the system shown in Figure **1A** suitable for use in implementing some embodiments of the present disclosure.

[0010] Figure **1C** illustrates a conceptual diagram of temporal consistency constraints, in accordance with an embodiment.

[0011] Figure **1D** illustrates a flowchart of a method for applying self-supervised adaptation to the system shown in Figure **1A** suitable for use in implementing some embodiments of the present disclosure.

[0012] Figure **2A** illustrates a block diagram of an example training configuration for the 3D object construction system shown in Figure **1A** suitable for use in implementing some embodiments of the present disclosure.

[0013] Figure **2B** illustrates a conceptual diagram of using temporal invariance to encourage parts correspondence, in accordance with an embodiment.

[0014] Figure **2C** illustrates a conceptual diagram of training using annotation re-projection suitable for use in implementing some embodiments of the present disclosure.

[0015] Figure **2D** illustrates a flowchart of a method for training the 3D object construction system shown in Figure **1A** suitable for use in implementing some embodiments of the present disclosure.

[0016] Figure **3** illustrates images and reconstructed objects, in accordance with an embodiment.

[0017] Figure **4** illustrates an example parallel processing unit suitable for use in implementing some embodiments of the present disclosure.

[0018] Figure **5A** is a conceptual diagram of a processing system implemented using the PPU of Figure **4** suitable for use in implementing some embodiments of the present disclosure.

[0019] Figure **5B** illustrates an exemplary system in which the various architecture and/or functionality of the various previous embodiments may be implemented.

[0020] Figure **6A** is a conceptual diagram of a graphics processing pipeline implemented by the PPU of Figure **4** suitable for use in implementing some embodiments of the present disclosure.

[0021] Figure **6B** illustrates an exemplary game streaming system suitable for use in implementing some embodiments of the present disclosure.

DETAILED DESCRIPTION

[0022] The task of 3D reconstruction entails the simultaneous recovery of the 3D shape, texture, and camera pose of objects from 2D images. The task is highly ill-posed due to the inherent ambiguity of correctly estimating both the shape and camera pose together. However, a 3D object construction neural network system may learn to predict 3D representations of the object from video.

[0023] In an embodiment, the 3D object construction neural network system is trained to reconstruct temporally consistent 3D meshes of deformable object instances from videos. In an embodiment, the videos include real animals in natural environments. Prior to inferencing, the neural network system is trained to jointly predict the shape, texture, and camera pose of an image for category-specific 3D reconstruction using a collection of single-view images of the same category. A first example category may include, but is not limited to, birds (including ducks). A second example category may be horses. In general, a category includes animals that have a similar structure, such as animals within a single species. The neural network may be trained without requiring an annotated 3D mesh, 2D keypoints, or camera pose for each video frame.

[0024] Then, at inference time, the neural network system is adapted over time using self-supervised regularization terms that exploit temporal consistency of an object instance to enforce that all reconstructed meshes for the object share a common texture map, a base (identity) shape, as well as parts. As a result of the adaptive refinement, the neural network system recovers temporally consistent and reliable 3D structures from videos of non-rigid objects including those of animals captured in the wild – a challenging task that is rarely addressed.

[0025] Figure 1A illustrates a block diagram of a 3D object construction system 100, in accordance with an embodiment. It should be understood that this and other arrangements described herein are set forth only as examples. Other arrangements and elements (e.g., machines, interfaces, functions, orders, groupings of functions, etc.) may be used in addition to or instead of those shown, and some elements may be omitted altogether. Further, many of the elements described herein are functional entities that may be implemented as discrete or distributed components or in conjunction with other components, and in any suitable

combination and location. Various functions described herein as being performed by entities may be carried out by hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory. Furthermore, persons of ordinary skill in the art will understand that any system that performs the operations of the 3D object construction system **100** is within the scope and spirit of embodiments of the present invention.

[0026] The 3D object construction system **100** includes a neural network model comprising at least an encoder **105**, shape decoder **115**, and motion decoder **120**. The encoder **105** extracts features **110** from each frame (e.g., image) in the video. An input image **102** including an object and a predicted 3D object **104**, output by the 3D object construction system **100**, are illustrated in Figure **1A**. The features **110** are then processed by a variety of decoders to predict an identity shape, motion (offsets), texture, and camera (not shown). In an embodiment, the shape decoder **115** outputs an identity shape **116** that represents basic shapes of the same category (e.g., duck, flying bird, fat bird, standing bird, etc.). In an embodiment, the identity shape **116** is defined by a 3D mesh of vertices that define faces of the mesh surface. In an embodiment, the meshes for a particular category are deformed from a predefined sphere and have the same number of vertices/faces. The motion decoder **120** predicts offsets, ∇V relative to each vertex in the identity shape **116**. For each video frame, the offsets define shape deformations that are applied to the identity shape **116** and, over time, appear as movement.

[0027] In contrast with conventional 3D reconstruction techniques, the predicted shapes are not limited to those that are symmetric. An assumption of symmetry does not hold for most non-rigid animals, e.g., birds tilting their heads, or walking horses, etc. This is particularly important for recovering dynamic meshes in sequences, e.g., when a bird rotates its head, the 3D shape is no longer mirror symmetric. Therefore, the assumption of symmetry may be removed and the constructed meshes are allowed to fit more complex, non-rigid poses. Simply removing the symmetry assumption for the predicted vertex offsets leads to excessive freedom in shape deformation. To resolve undesirable deformations, a set of N_b shape bases $\{V_i\}_{i=1}^{N_b}$ is learned by the shape decoder **115** based on the features **110**. The base or identity shape **116** is computed as a weighted combination of the shape bases, denoted as the base shape V_{base} . Compared to a single mesh template, the base shape V_{base} is more powerful in capturing the object's identity and

relieves the shape decoder **115** from predicting large motion deformation, e.g., of deforming a standing bird template to a flying bird.

[0028] Different sets of shape bases are learned during training by clustering constructed meshes, where the meshes in each set share a similar shape and the base shape is their mean shape. Learned coefficients $\{\beta_i\}_{i=1}^{N_b}$ are predicted by the shape decoder **115** and used to combine corresponding meshes in the set of shape bases to produce the identity shape **116**. The identity shape **116** may be computed as:

$$V_{\text{base}} = \sum_{i=1}^{N_b} \beta_i V_i. \quad \text{Eq. (1)}$$

[0029] The motion decoder **120** predicts offsets ∇V relative to each vertex in the identity shape **116**. The offsets **118** encode the object’s asymmetric non-rigid motion, defining deformations for each vertex in the identity shape **116**.

$$V = V_{\text{base}} + \nabla V. \quad \text{Eq. (2)}$$

The offsets **118** are applied to the vertices of the identity shape **116** by 3D mesh construction unit **112** to construct a predicted 3D shape representation **108** (e.g., wireframe or mesh) of the object. The predicted 3D shape representation **108** is shown as a checkerboard in Figure **1A** for visualization purposes, where different faces defined by the vertices in the mesh are colored either black or white.

[0030] The texture decoder **125** receives the features **110** and predicts a texture image **106** for each frame of the video. The texture decoder **125** predicts a texture flow $I_{\text{flow}} \in \mathbb{R}^{H_{\text{uv}} \times W_{\text{uv}} \times 2}$ for each image based on the features **110**. The texture flow maps pixels from the image to UV texture space to produce a texture image **106**. A predefined UV mapping function may then be used by 3D mesh texture unit **122** to map the texture image **106** from the texture space to the 3D shape representation **108**. Applying the texture image **106** to the 3D shape representation **108** produces a 3D object **104** (e.g., textured mesh). The 3D object **104** is represented with $|V|$ vertices ($V \in \mathbb{R}^{|V| \times 3}$), $|F|$ faces ($F \in \mathbb{R}^{|F| \times 3}$), and the UV texture image $I_{\text{uv}} \in \mathbb{R}^{H_{\text{uv}} \times W_{\text{uv}} \times 3}$ of height H_{uv} and width W_{uv} . The UV texture space provides a parameterization that is invariant to object deformation. Therefore, over time, the image texture for the images in the video should be constant or invariant to shape deformation.

[0031] By enforcing the predicted values for the texture images **106** to be consistent in the UV texture space across different frames of a video, the neural network model can be regularized to generate coherent reconstructions over time during inferencing. The temporal invariance may be used as self-supervised signals to tune the neural network model. In an embodiment, during inferencing, the 3D object construction system **100** is adapted using self-supervised regularization based on shape invariance and texture invariance. Specifically, an object in a video maintains a consistent shape and consistent texture across multiple frames. The self-supervised adaptation technique exploits the temporal consistency of the object to construct the dynamic 3D object from an unlabeled video. As described further herein, texture image, identity shape, and part correspondence invariance constraints may be applied to train and/or fine-tune the neural network model.

[0032] More illustrative information will now be set forth regarding various optional architectures and features with which the foregoing framework may be implemented, per the desires of the user. It should be strongly noted that the following information is set forth for illustrative purposes and should not be construed as limiting in any manner. Any of the following features may be optionally incorporated with or without the exclusion of other features described.

[0033] Figure **1B** illustrates flowchart of a method **130** for constructing a 3D representation of an object using the 3D object reconstruction system **100** shown in Figure **1A**, in accordance with an embodiment. Each block of method **130**, described herein, comprises a computing process that may be performed using any combination of hardware, firmware, and/or software. For instance, various functions may be carried out by a processor executing instructions stored in memory. The method **130** may also be embodied as computer-usable instructions stored on computer storage media. The method **130** may be provided by a standalone application, a service or hosted service (standalone or in combination with another hosted service), or a plug-in to another product, to name a few. In addition, method **130** is described, by way of example, with respect to the system of Figure **1A**. However, this method may additionally or alternatively be executed by any one system, or any combination of systems, including, but not limited to, those described herein. Furthermore, persons of ordinary skill in the art will

understand that any system that performs method **130** is within the scope and spirit of embodiments of the present invention.

[0034] At step **135**, the 3D object construction system **100** receives a video including images of the object captured from a camera pose. In an embodiment, the video is unlabeled. In an embodiment, the object is a non-rigid animal. In an embodiment, the video is captured in the “wild”.

[0035] At step **140**, the 3D object construction system **100** predicts the 3D shape representation of the object for a first image of the images based on a set of learned shape bases. In an embodiment, the 3D object construction system **100** predicts an identity shape from a set of learned shapes. In an embodiment, the identity shape is computed using Equation (1). In an embodiment, the identity shape is computed as a sum of component shapes included in the set of learned shape bases and each component shape is corresponding scaled by a coefficient generated by the neural network model. In an embodiment, shape offsets (e.g., non-rigid motion deformations) are computed and applied to the vertices of the identity shape to predict the 3D shape representation of the object. In an embodiment, the 3D shape representation is a mesh of vertices that define faces.

[0036] At step **145**, the 3D object construction system **100** predicts a texture flow for the first image. At step **147**, the 3D object construction system **100** maps pixels from the first image to a texture space according to the texture flow to produce a texture image corresponding to the first image. In an embodiment, for a particular image, the texture image is transferred onto the 3D shape representation to produce the 3D object corresponding to the object in the first image. In an embodiment, steps **140**, **145**, and **147** are repeated for each image in the video. In an embodiment, the 3D construction system **100** also predicts a camera pose based on features extracted from the video. When the 3D object is rendered according to the camera pose for each frame of video, the rendered object appears as the object in the frame.

[0037] Figure **1C** illustrates a conceptual diagram of temporal consistency constraints, in accordance with an embodiment. A neural network model **150** comprises the encoder **105**, shape decoder **115**, and texture decoder **125**. In an embodiment, to enforce texture invariance, the reconstructed texture images for an arbitrary pair of frames (e.g., images) in a video sequence are swapped. As shown in Figure **1C**, frames **132** and **136** are processed by the neural network

model **150** to predict the respective identity shapes **142** and **146**. The identity shapes **142** and **146** shown in Figure **1C** are each posed according to a respective camera pose, θ of the respective frames, where $\theta \in \mathbb{R}^7$. In an embodiment, the camera pose represents a perspective transformation. Except for the different camera poses, the identity shapes **142** and **146** are the same. The offsets predicted by the neural network model **150** for the frames **132** and **136** are applied to the respective identity shapes **142** and **146**, to produce 3D shape representations **152** and **156**.

[0038] The shape model is represented by an identity or base shape V_{base} and an offset or deformation term ΔV , in which the identity shape V_{base} intuitively corresponds to the “identity” of the instance, e.g., a duck, or a flying bird, etc. During online adaptation, the neural network model **150** is trained to predict consistent V_{base} to preserve the identity shape over time, via a swapping loss function. Given two randomly sampled frames, I^i and I^j , the identity shapes V_{base}^i and V_{base}^j are swapped and deformed using the original offsets ΔV^i and ΔV^j as:

$$L_s = \text{niou}(\mathcal{R}(V_{\text{base}}^j + \Delta V^i, \theta^i), S^i) + \text{niou}(\mathcal{R}(V_{\text{base}}^i + \Delta V^j, \theta^j), S^j), \quad \text{Eq. (3)}$$

where θ^i and θ^j are the camera poses, ΔV^i and ΔV^j are the motion deformations, and S^i and S^j are the object silhouettes (masks) of frame i and j , respectively. $\mathcal{R}(\cdot)$ denotes a general projection. In an embodiment, $\mathcal{R}(\cdot)$ represents a differentiable renderer to render the 3D representation or 3D object to a 2D silhouette as $\mathcal{R}(V, \theta)$. In an embodiment, $\mathcal{R}(\cdot)$ represents a differentiable renderer to render a texture mesh to an RGB image as $\mathcal{R}(V, \theta, I_{\text{uv}})$ (note that mesh faces F are omitted for conciseness). In an embodiment, $\mathcal{R}(\cdot)$ represents a projection of a 3D point v to the image space as $\mathcal{R}(v, \theta)$. The $\text{niou}(\cdot, \cdot)$ denotes the negative intersection over union (IoU) objective. In an embodiment, the foreground masks used for the silhouette are obtained by a segmentation model trained with the ground truth foreground masks. The loss function L_s may be used to enforce consistency between the identity shapes **142** and **146**.

[0039] The neural network model **150** also predicts the texture images **154** and **158** for the frames **132** and **136**, respectively. Based on the observation that object texture mapped to the UV space should be invariant to shape deformation and stay constant over time, a texture invariance constraint may be used to encourage consistent texture reconstruction from all frames. However, naively aggregating the UV texture maps from all the frames may lead to a blurry

video-level texture image. Instead, like the shape identity, texture consistency may be enforced between random pairs of frames, via a swap loss. Given two randomly sampled frames, I^i and I^j , the texture images I_{uv}^i and I_{uv}^j are swapped and combined with the original mesh reconstructions V^i and V^j as:

$$L_t = \text{dist}(\mathcal{R}(V^i, \theta^i, I_{uv}^j) \odot S^i, I^i \odot S^i) + \text{dist}(\mathcal{R}(V^j, \theta^j, I_{uv}^i) \odot S^j, I^j \odot S^j), \quad \text{Eq. (4)}$$

where $\text{dist}(\cdot, \cdot)$ is a perceptual metric. The swapping technique enforces consistency of the texture images across the frames, thereby improving the accuracy of the predicted texture images **106**, 3D shape representations **108**, and 3D objects **104** in Figure **1A**. For example, the loss function L_t may be used to enforce consistency between the texture images **154** and **158**. During inferencing, the neural network model **150** may be fine-tuned on a particular video with the invariance constraints enforced by Equations (3) and (4). In an embodiment, the neural network model **150** is fine-tuned using self-supervision. Fine-tuning may improve the performance of the neural network model **150** when domain differences in video quality, such as lighting conditions, etc. cause inconsistent 3D mesh reconstruction as each frame of a video is processed independently.

[0040] Figure **1D** illustrates a flowchart of a method **148** for applying self-supervised adaptation to the system shown in Figure **1A**, in accordance with an embodiment. The neural network model **150** is refined by taking advantage of the redundancy in temporal sequences as a form of self-supervision in order to improve the construction of dynamic non-rigid objects. Although method **148** is described in the context of a neural network model, the method **148** may also be performed by a program, custom circuitry, or by a combination of custom circuitry and a program. For example, the method **148** may be executed by a GPU, CPU, or any processor capable of implementing the neural network model. Furthermore, persons of ordinary skill in the art will understand that any system that performs method **148** is within the scope and spirit of embodiments of the present invention.

[0041] At step **155**, the neural network model **150** receives a sequence of frames for a video including an object. At step **160**, silhouettes and camera poses are obtained. In an embodiment, the silhouettes are ground truth data corresponding to the video. In an embodiment, the camera poses are predicted by the neural network model **150**. Steps **165** and

170 may be performed in parallel or in sequence. At step **165**, texture images are predicted for each frame by the neural network model **150**. At step **170**, an identity shape and offsets are predicted for each frame by the neural network model **150**.

[0042] At step **175**, a loss function is computed based on texture invariance and shape identity invariance. In an embodiment, the loss function is a swapping loss function. In an embodiment, the loss function is a combination of Equations (3) and (4).

[0043] At step **180**, if the error is reduced according to the loss function, then at step **190**, refinement of the neural network model **150** is completed. Otherwise, at step **185**, parameters of the neural network model **150** are updated by back propagating the loss function before the method **148** returns to step **155**. In an embodiment, parameters of one or more of the shape decoder **115**, motion decoder **120**, and texture decoder **125** are updated.

[0044] In an embodiment, the texture image predicted for a first image is transferred to a second 3D shape representation predicted for a second image of the images in a video to produce a first 3D object. The first 3D object is projected according to a first camera pose associated with the first image to produce a first projected 3D object. A second texture image predicted for the second image is transferred to the 3D shape representation predicted for the first image to produce a second 3D object. The second 3D object is projected, according to a second camera pose associated with the second image, to produce a second projected 3D object and parameters of the neural network model **150** are updated to encourage consistency between the first projected 3D object and the second projected 3D object. In an embodiment, differences between overlapped parts of the first and second image are reduced during training.

[0045] In an embodiment, first non-rigid motion deformations predicted for the first image are applied to a first identity shape predicted for a second image of the images to produce a first 3D shape representation. The first 3D shape representation is projected according to a first camera pose associated with the first image to produce a first projected 3D object. Second non-rigid motion deformations predicted for the second image are applied to a second identity shape predicted for the first image to produce a second 3D shape representation. The second 3D shape representation is projected according to a second camera pose associated with the second image to produce a second projected 3D object and parameters of the neural network model **150**

are updated to encourage consistency between the first projected 3D object and the second projected 3D object.

[0046] Figure 2A illustrates a block diagram of a training configuration for the 3D object construction system **100** shown in Figure 1A, in accordance with an embodiment. In addition to the blocks shown in Figure 1A, the training framework includes a camera pose unit **225**, a differentiable renderer **222**, and a loss function **215**. In an embodiment, the camera pose unit **225** is also included in the neural network model **150** and/or the 3D object construction system **100**. The camera pose unit **225** receives the features and predicts a camera pose (position and orientation), θ of the identity shape based on the features. The shape decoder **115**, motion decoder **120**, texture decoder **125**, and camera pose unit **225** may be configured to jointly predict the identity shape, offsets, texture image, and camera pose.

[0047] The differentiable renderer **222** receives the 3D mesh and texture image and renders the 3D mesh according to the predicted camera pose provided by the camera pose unit **225**. In an embodiment, the texture image is transferred onto the 3D shape representation by the 3D mesh texture unit **122** to construct the 3D object. The textured 3D object is then projected by the differentiable renderer **222** according to the camera pose to produce a rendered image.

[0048] Rendered images may be compared with the video frames and/or annotations during training. The loss function **215** may adjust parameters of the shape decoder **115**, motion decoder **120**, texture decoder **125**, and/or camera pose unit **225** to reduce differences between the rendered images and video frames. In some embodiments, the loss function **215** receives ground truth data, such as object silhouettes and/or keypoint annotations, that are used to reduce differences between the rendered images and video frames. In an embodiment, training is accomplished in a self-supervised manner using object silhouettes extracted from category-specific images. The training configuration may be used to train the 3D object construction system **100** using both individual input images and videos. In an embodiment, training may be performed using a variety of techniques, e.g., supervised, self-supervised, and semi-supervised using individual images and/or videos. For example, the 3D object construction system **100** may be trained to learn sets of shape bases using semi-supervised techniques with single-view input images, followed by fine-tuning via self-supervised techniques using an unlabeled video of a specific object.

[0049] Conventional techniques use annotated 2D object keypoints and category-level template shapes or silhouettes for training. However, scaling up learning with 2D annotations to hundreds of thousands of images is non-trivial and may also limit the generalization ability of a trained neural network model for new domains. For example, a 3D construction neural network model that is conventionally trained on single-view images typically produces unstable and erratic predictions for video data. This is unsurprising, due to perturbations over time. Therefore, the temporal signal in videos should be used to provide an advantage instead of a disadvantage.

[0050] A balance may be achieved between model generalization and specialization. In an embodiment, an image-based neural network model is trained on a set of images, while at test time, the neural network model within the 3D object construction system **100** is adapted or fine-tuned to an input video including a particular identity object. During test-time training no labels are provided for the video. Therefore, self-supervised objectives are introduced that can continuously improve the neural network model. As previously described, the UV texture space provides a parameterization that is invariant to object deformation. Object parts of an instance of an object should be constant when mapped from 2D via the predicted texture flow. Therefore, in addition to encouraging temporal consistency between the texture images and identity shapes predicted for different frames, temporal consistency may be encouraged between object parts in UV texture space. In particular, invariance loss may be used to train the camera pose unit **225** to predict the camera position θ . Using the constraint of temporal consistency, the recovered identity shape and camera pose may be stabilized considerably and adapted based on the video that is processed during fine-tuning and/or test time.

[0051] Figure **2B** illustrates a conceptual diagram of using temporal invariance to enforce parts correspondence, in accordance with an embodiment. The video includes a bird object that the 3D object construction system **100** may be trained to construct. Conventional techniques, such as unsupervised video correspondence (UVC) may be used to automatically apply (random) patterns to parts of the object in input video frames to generate a propagated part map for each frame. The UVC model learns an affinity matrix that captures pixel-level correspondences among video frames. The UVC model can be used to propagate any annotation (e.g., segmentation labels, keypoints, part labels, etc.), from an annotated keyframe to the other unannotated frames. Part correspondence is generated within a video clip by "painting" a group

of random regions on the object on the first frame and propagating the part painting to the rest of the video using the UVC model. Any specific part of the object, such as a wing may be painted as a single region or as multiple regions. In other words, the painting does not provide a semantic definition. As shown in Figure **2B**, the painting appears as differently colored vertical stripes within a silhouette of the object visible in a propagated part map **232**.

[0052] In an embodiment, to obtain accurate part propagation of object parts by the UVC model, two strategies may be employed. Firstly, parameters in the 3D object construction system **100** may be fine-tuned on sliding windows instead of all video frames. Each sliding window may include $N_w=50$ consecutive frames and the sliding stride is set to $N_s=10$. The 3D object construction system **100** may be tuned for $N_t=40$ iterations with frames in each sliding window. Secondly, instead of “painting” random parts onto the first frame and propagating the painted parts to rest of the frames sequentially in a window, random parts may be painted onto the *middle* frame (i.e. the $N_w/2^{\text{th}}$ frame) in the window and the painted parts may be propagated backward to the first frame as well as forward to the last frame in the window. The strategy may improve the propagation quality by decreasing the propagation range to half of the window size. Within each sliding window, consistency of the UV texture images, part UV maps, as well as identity shapes is encouraged for all of the frames.

[0053] A sequence of propagated part maps associated with video frames, such as the propagated part map **232** associated with the frame **230** is processed by the 3D object construction system **100** in a test configuration. The processing of the propagated part map **232** is described, however, additional propagated part maps shown in Figure **2B** or additional frames in the sequence may be processed in a similar manner to produce additional rendered images. In an embodiment, the part map is propagated across the object in a number of the frames to produce propagated part maps. In an embodiment the number of frames is included in a sliding window.

[0054] The propagated part map for each frame is mapped to the UV texture space with the predicted texture flow **234** to produce part UV map **236**. In an embodiment, the propagated part maps are mapped into the texture space according to corresponding texture flows predicted for the frames to produce part maps in the texture space. In an embodiment, the part maps are aggregated to produce a video level part UV map **235**. By aggregating the part UV maps, i.e.,

averaging, noise is minimized in each individual part UV map. The video-level part UV map **235** for the object depicted in the video that will be constructed in 3D is shared by all frames in the video. Thus, for each frame, the video-level part UV map **235** is wrapped onto the predicted 3D shape representation **237**. Aggregation of the predicted parts in the UV texture space facilitates learning the camera pose by the 3D object construction system **100**.

[0055] In an embodiment, 3D shape representations predicted for the frames are rendered according to the associated camera poses (not shown) and the video-level part map is transferred (e.g., wrapped) onto each one of the 3D shape representations to produce rendered images. The wrapped 3D shape representation **237** may be rendered by the differentiable renderer **222** according to the predicted camera to produce rendered image **238**.

[0056] Discrepancies between the parts rendered back to the 2D space and the part propagations, for each frame, are penalized. In an embodiment, the loss function **215** may be configured to update parameters of the neural network model **150** based on the discrepancies. In an embodiment, the parameters are updated to encourage consistency between the rendered images and the propagated part maps.

[0057] As the propagated part maps are usually temporally smooth and continuous, the loss implicitly regularizes the 3D object construction system **100** to predict coherent camera pose and object shape over time. In an embodiment, instead of minimizing the discrepancy between the rendered part map (e.g., rendered image, such as the rendered image **238**) and the propagated part map of a frame, it may be more robust to penalize the geometric distance between the projections of vertices assigned to each part with 2D points sampled from the corresponding part. A Chamfer loss may be computed as:

$$L_c = \sum_{j=1}^{N_f} \sum_{i=1}^{N_p} \frac{1}{|V_i^j|} \text{Chamfer}(\mathcal{R}(V_i^j, \theta^j), Y_i^j), \quad \text{Eq. (5)}$$

where N_f is the number of frames in the video, $N_p=6$ is the number of parts, and V_i^j are vertices assigned to part i . The Chamfer distance is used to compute the loss because the vertex projections $\mathcal{R}(V_i^j, \theta^j)$ are not strictly one-to-one corresponded to the sampled 2D points Y_i^j .

[0058] The input samples of the propagated part maps are compared with the projected samples of the images rendered based on the predicted video-level part UV map, 3D

representation, and camera pose to calculate the Chamfer loss. The Chamfer loss reduces errors from the camera pose estimation. Alternatively, more samples (e.g., pixels) within the colored parts of the propagated part map **232** may be used as the ground truth and the predicted 3D shape representation **237** may be wrapped with the video-level part UV map **235** and rendered according to the camera pose to produce a rendered image **238** for comparison with the ground truth colored samples.

[0059] In an embodiment, the colored parts are used to supervise the estimated camera pose without rendering images. The input video frames may be sampled within each part of the part propagations and then compared with the predicted samples on the 3D representations (e.g., meshes) projected according to the predicted camera pose (no rendering).

[0060] Another technique for training uses propagation of the ground truth annotations, such as keypoints, from the input images through the predicted texture images to the rendered images. A loss may be computed based on the ground truth annotations and the rendered annotations. The ground truth annotations enable correspondences to be established across different instances in a set of shape bases. For example, a beak or wingtip is labeled as a keypoint in different images of birds and different birds with similar shapes are clustered together to form a set of shape bases.

[0061] Figure **2C** illustrates a conceptual diagram of training using annotation re-projection, in accordance with an embodiment. Although the technique is described for annotations that are keypoints, those skilled in the art will recognize that other types of annotations may be used with the technique. When the 3D object construction system **100** is trained using weak supervision, 2D keypoints may be provided as ground truth annotations that semantically associate different instances of the object. For example, an annotated frame **240** includes multiple keypoints, such as a tail keypoint **248** at the tip of the bird's tail.

[0062] When the 2D keypoints are projected onto the 3D representation (e.g., mesh surface), the same semantic keypoint for different object instances should be matched to the same face on the mesh surface. Conventionally, to model the mapping between the 3D mesh surface and the 2D keypoints, an affinity matrix is learned that describes the probability of each 2D keypoint

mapping to each vertex on the 3D representation. The probability map is a heatmap, such as heatmap **242**.

[0063] The affinity matrix is shared among all instances and is independent of individual shape variations. The conventional approach is sub-optimal because: (i) Mesh vertices are a subset of discrete points on a continuous mesh surface and so their weighted combination defined by the affinity matrix may not lie on the mesh surface, leading to inaccurate mappings of 2D keypoints. (ii) The mapping from the image space to the mesh surface is described by the affinity matrix. In contrast, because the mapping of image space the mesh surface is already modeled by the texture flow, it is potentially redundant to independently learn both the affinity matrix and the texture flow.

[0064] Therefore, the texture flow is re-utilized to map 2D keypoints from each image to the mesh surface. For example, the 2D tail keypoint **248** is mapped to a 3D tail keypoint **246** on the 3D representation. First, each 2D keypoint in an annotated frame is mapped to the UV texture space that is independent of shape deformation. For example, the keypoints in the annotated frame **240** may be mapped to the UV texture space according to the texture flow, I_{flow} that is predicted for the annotated frame **240** to generate an annotation map in the texture space, keypoint UV map **244**. Ideally, each semantic keypoint from different instances of the object should map to the same point in the UV space. In practice, this may not hold due to inaccurate texture flow prediction. To accurately map each keypoint to the UV space, a canonical keypoint UV map **245** may be computed by: (i) mapping the keypoint heat map for each instance to the UV space via the predicted texture flow, and (ii) aggregating the keypoint UV maps across all instances to eliminate outliers caused by incorrect texture flow prediction. The keypoint UV maps are aggregated (e.g., averaged) to produce the canonical keypoint UV map **245**.

[0065] The canonical keypoint UV map **245** is then transferred to 3D shape representations predicted for the frames to produce annotated 3D shape representations. The annotated 3D shape representations may then be projected, according to associated camera poses, to produce projected annotations for the frames. In an embodiment, the projection comprises rendering. In an embodiment, the keypoint re-projection is done by (i) warping the canonical keypoint UV map to each individual predicted mesh surface to produce a 3D keypoint; (ii) projecting the 3D keypoints K_{3D}^i back to the 2D space via the predicted camera pose to produce re-projected 2D

keypoints; (iii) comparing the re-projected 2D keypoints against the ground truth keypoints in 2D, K_{2D}^i . For example, the canonical keypoint UV map **245** is mapped to the 3D representation to produce 3D keypoint **246**, including the tail keypoint **250**. The 3D keypoints **246** are re-projected according to the predicted camera pose and compared with the ground truth keypoints in the annotated frame **240**. In an embodiment, parameters of the neural network model **150** are updated to reduce differences and encourage consistency between the projected keypoints and the ground truth keypoints.

[0066] Given the 3D correspondence (denoted as K_{3D}^i of each 2D semantic keypoint K_{2D}^i , a keypoint re-projection loss enforces the projection of the former to be consistent with the latter by:

$$L_{kp} = \frac{1}{N_k} \sum_{i=1}^{N_k} \|\mathcal{R}(K_{3D}^i, \theta) - K_{2D}^i\|, \quad \text{Eq. (6)}$$

where N_k is the number of keypoints. Evaluation of the keypoint re-projection loss function implicitly reveals the correctness of both the predicted shape and camera pose for the mesh reconstruction algorithm, especially for objects that do not have 3D ground truth annotations. The re-projection of annotations enables weakly supervised training of the 3D object construction system **100**.

[0067] One bottleneck of conventional image-based 3D mesh reconstruction methods is that the predicted shapes are assumed to be symmetric. This assumption does not hold for most non-rigid animals, e.g., birds tilting their heads, or walking horses, etc. Therefore, the assumption of symmetry may be ignored and the reconstructed mesh representations may be allowed to fit more complex, non-rigid poses via an as-rigid-as-possible (ARAP) constraint. The ARAP constraint is an additional loss objective that may be used for self-supervised training of the 3D object construction system **100**. The identity shape is smooth by construction. However, application of the offsets predicted by the motion decoder **120** may produce discontinuities in the 3D mesh representation. The ARAP constraint is used to ensure that edge lengths of the 3D mesh are maintained even when the 3D mesh is rotated. ARAP is a self-supervised regularization that maintains rigidity and can be used on individual input images and video sequences.

[0068] Without any pose-related regularization, the predicted motion deformation ΔV often leads to erroneous random deformations and spikes on the surface of the 3D mesh, which do not

faithfully describe the motion of a non-rigid object. The ARAP constraint encourages rigidity of local transformations and the preservation of the local mesh structure. The ARAP constraint is formulated as an objective that ensures that the predicted shape V is a locally rigid transformation from the predicted base shape V_{base} by:

$$L_{\text{arap}}(V_{\text{base}}, V) = \sum_{i=1}^{|V|} \sum_{j \in \mathcal{N}(i)} w_{ij} \|(V^i - V^j) - R_i(V_{\text{base}}^i - V_{\text{base}}^j)\|, \quad \text{Eq. (7)}$$

where $\mathcal{N}(i)$ represents the neighboring vertices of a vertex i , w_{ij} and R_i , are the cotangent weight and the best approximating rotation matrix, respectively. As another constraint that does not require any labels, ARAP may be enforced during test-time training with video inputs, to improve shape prediction.

[0069] In an embodiment, parameters of the neural network model **150** are updated based on the ARAP loss function L_{arap} to reduce discontinuities in the predicted 3D shape representations. Non-rigid motion deformations of the 3D shape representations are predicted for the frames and applied to identity shapes predicted for the frames to produce the 3D shape representations of the object. The ARAP loss function may be evaluated based on rotated differences between the identity shapes and differences between the 3D shape representations.

[0070] In an embodiment, two image-based 3D object construction methods are used for training the neural network model **150** (i) a weakly supervised method (i.e., with object silhouettes and 2D annotations provided), and (ii) a self-supervised method where only object silhouettes are available. The image-based trained neural network model **150** is then adapted to videos. For example, in an embodiment, the trained neural network model **150** is adapted to videos that are in-the-wild bird and zebra videos.

[0071] Figure **2D** illustrates a flowchart of a method **255** for training the 3D object construction system **100** shown in Figure **1A**, in accordance with an embodiment. At step **260**, the 3D object construction system **100** is trained to learn a set of shape bases from single-view images. In an embodiment, objectives used individually or in combination for the single-view construction include (i) foreground mask loss: a negative intersection over union objective between rendered and ground truth silhouettes; (ii) foreground RGB texture loss: a perceptual metric between rendered and input RGB images; (iii) mesh smoothness: a Laplacian constraint to

encourage smooth mesh reconstruction; (iv) keypoint re-projection loss; and (v) the ARAP constraint.

[0072] At step **270**, the 3D object construction system **100** is trained using self-supervision for videos. Since it is feasible to predict a segmentation mask via a pretrained segmentation model, the predicted foreground mask may be used to compute the foreground mask loss. In an embodiment, objectives used individually or in combination for the for videos may include the foreground RGB texture loss, and the mesh smoothness objective. In an embodiment, the ARAP constraint may also be used.

[0073] At step **275**, the 3D object construction system **100** is fine-tuned for constructing a particular 3D object using the one or more of the invariance constraints for texture, identity shape, and part correspondence.

[0074] Figure **3** illustrates images and reconstructed objects, in accordance with an embodiment. The 2D bird object shown in images **300**, **310**, and **320** is constructed in 3D using the 3D object construction system **100** including the camera pose unit **225** trained on single-view images to produce the 3D representations and objects **305**, trained on single view images and videos, but without invariance constraints, to produce the 3D representations and objects **315**, and trained on single view images and videos with invariance constraints to produce the 3D representations and objects **325**. Note that the object shape, camera pose, and texture is less reliably predicted for the 3D representations and objects **315** compared with **325**. Similarly, the object shape, camera pose, and texture are less reliably predicted for the 3D representations and objects **305** compared with **315**. In sum, the 3D object construction algorithm recovers temporally consistent and reliable 3D structures from videos of non-rigid objects including those of animals captured in the wild.

[0075] The 3D object construction technique does not require pre-defined template object meshes, annotations of a 3D object, 2D annotations, or camera pose for the video frames. The video-based 3D object construction system **100** may be refined via self-supervised online adaptation for any incoming test video. First, a category-specific 3D construction neural network model **150** is learned from a collection of single-view images of the same category. The 3D object construction system **100** jointly predicts the shape, texture, and camera pose of an object in an image. Then, at inference time, the neural network model is fine-tuned over time by an

object-specific test video using self-supervised regularization terms that exploit temporal consistency of an object instance to enforce that all reconstructed meshes share a common texture map, a base shape, as well as parts.

[0076] The 3D object construction technique may be used for content creation, such as generation of 3D characters for games, movies, and 3D printing. Because the 3D characters are generated from video, the content may also include motion of the character, as predicted based on the video. Compared with conventional solutions, the 3D object construction technique does not rely on a pre-defined parametric mesh (e.g., human object defined by a fixed number of vertices). The 3D object construction system also generalizes well, particularly for non-rigid objects.

Parallel Processing Architecture

[0077] Figure 4 illustrates a parallel processing unit (PPU) **400**, in accordance with an embodiment. The PPU **400** may be used to implement the 3D object construction system **100**. The PPU **400** may be used to implement one or more of the encoder **105**, shape decoder **115**, motion decoder **120**, texture decoder **125**, 3D mesh construction unit **130**, and 3D mesh texture unit **122** within the server/client system **100**.

[0078] In an embodiment, the PPU **400** is a multi-threaded processor that is implemented on one or more integrated circuit devices. The PPU **400** is a latency hiding architecture designed to process many threads in parallel. A thread (e.g., a thread of execution) is an instantiation of a set of instructions configured to be executed by the PPU **400**. In an embodiment, the PPU **400** is a graphics processing unit (GPU) configured to implement a graphics rendering pipeline for processing three-dimensional (3D) graphics data in order to generate two-dimensional (2D) image data for display on a display device. In other embodiments, the PPU **400** may be utilized for performing general-purpose computations. While one exemplary parallel processor is provided herein for illustrative purposes, it should be strongly noted that such processor is set forth for illustrative purposes only, and that any processor may be employed to supplement and/or substitute for the same.

[0079] One or more PPUs **400** may be configured to accelerate thousands of High Performance Computing (HPC), data center, cloud computing, and machine learning

applications. The PPU **400** may be configured to accelerate numerous deep learning systems and applications including autonomous vehicle platforms, deep learning, high-accuracy speech, image, and text recognition systems, intelligent video analytics, molecular simulations, drug discovery, disease diagnosis, weather forecasting, big data analytics, astronomy, molecular dynamics simulation, financial modeling, robotics, factory automation, real-time language translation, online search optimizations, and personalized user recommendations, and the like.

[0080] As shown in Figure **4**, the PPU **400** includes an Input/Output (I/O) unit **405**, a front end unit **415**, a scheduler unit **420**, a work distribution unit **425**, a hub **430**, a crossbar (Xbar) **470**, one or more general processing clusters (GPCs) **450**, and one or more memory partition units **480**. The PPU **400** may be connected to a host processor or other PPUs **400** via one or more high-speed NVLink **410** interconnect. The PPU **400** may be connected to a host processor or other peripheral devices via an interconnect **402**. The PPU **400** may also be connected to a local memory **404** comprising a number of memory devices. In an embodiment, the local memory may comprise a number of dynamic random access memory (DRAM) devices. The DRAM devices may be configured as a high-bandwidth memory (HBM) subsystem, with multiple DRAM dies stacked within each device.

[0081] The NVLink **410** interconnect enables systems to scale and include one or more PPUs **400** combined with one or more CPUs, supports cache coherence between the PPUs **400** and CPUs, and CPU mastering. Data and/or commands may be transmitted by the NVLink **410** through the hub **430** to/from other units of the PPU **400** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). The NVLink **410** is described in more detail in conjunction with Figure **5B**.

[0082] The I/O unit **405** is configured to transmit and receive communications (e.g., commands, data, etc.) from a host processor (not shown) over the interconnect **402**. The I/O unit **405** may communicate with the host processor directly via the interconnect **402** or through one or more intermediate devices such as a memory bridge. In an embodiment, the I/O unit **405** may communicate with one or more other processors, such as one or more the PPUs **400** via the interconnect **402**. In an embodiment, the I/O unit **405** implements a Peripheral Component Interconnect Express (PCIe) interface for communications over a PCIe bus and the interconnect

402 is a PCIe bus. In alternative embodiments, the I/O unit **405** may implement other types of well-known interfaces for communicating with external devices.

[0083] The I/O unit **405** decodes packets received via the interconnect **402**. In an embodiment, the packets represent commands configured to cause the PPU **400** to perform various operations. The I/O unit **405** transmits the decoded commands to various other units of the PPU **400** as the commands may specify. For example, some commands may be transmitted to the front end unit **415**. Other commands may be transmitted to the hub **430** or other units of the PPU **400** such as one or more copy engines, a video encoder, a video decoder, a power management unit, etc. (not explicitly shown). In other words, the I/O unit **405** is configured to route communications between and among the various logical units of the PPU **400**.

[0084] In an embodiment, a program executed by the host processor encodes a command stream in a buffer that provides workloads to the PPU **400** for processing. A workload may comprise several instructions and data to be processed by those instructions. The buffer is a region in a memory that is accessible (e.g., read/write) by both the host processor and the PPU **400**. For example, the I/O unit **405** may be configured to access the buffer in a system memory connected to the interconnect **402** via memory requests transmitted over the interconnect **402**. In an embodiment, the host processor writes the command stream to the buffer and then transmits a pointer to the start of the command stream to the PPU **400**. The front end unit **415** receives pointers to one or more command streams. The front end unit **415** manages the one or more streams, reading commands from the streams and forwarding commands to the various units of the PPU **400**.

[0085] The front end unit **415** is coupled to a scheduler unit **420** that configures the various GPCs **450** to process tasks defined by the one or more streams. The scheduler unit **420** is configured to track state information related to the various tasks managed by the scheduler unit **420**. The state may indicate which GPC **450** a task is assigned to, whether the task is active or inactive, a priority level associated with the task, and so forth. The scheduler unit **420** manages the execution of a plurality of tasks on the one or more GPCs **450**.

[0086] The scheduler unit **420** is coupled to a work distribution unit **425** that is configured to dispatch tasks for execution on the GPCs **450**. The work distribution unit **425** may track a number of scheduled tasks received from the scheduler unit **420**. In an embodiment, the work

distribution unit **425** manages a pending task pool and an active task pool for each of the GPCs **450**. As a GPC **450** finishes the execution of a task, that task is evicted from the active task pool for the GPC **450** and one of the other tasks from the pending task pool is selected and scheduled for execution on the GPC **450**. If an active task has been idle on the GPC **450**, such as while waiting for a data dependency to be resolved, then the active task may be evicted from the GPC **450** and returned to the pending task pool while another task in the pending task pool is selected and scheduled for execution on the GPC **450**.

[0087] In an embodiment, a host processor executes a driver kernel that implements an application programming interface (API) that enables one or more applications executing on the host processor to schedule operations for execution on the PPU **400**. In an embodiment, multiple compute applications are simultaneously executed by the PPU **400** and the PPU **400** provides isolation, quality of service (QoS), and independent address spaces for the multiple compute applications. An application may generate instructions (e.g., API calls) that cause the driver kernel to generate one or more tasks for execution by the PPU **400**. The driver kernel outputs tasks to one or more streams being processed by the PPU **400**. Each task may comprise one or more groups of related threads, referred to herein as a warp. In an embodiment, a warp comprises 32 related threads that may be executed in parallel. Cooperating threads may refer to a plurality of threads including instructions to perform the task and that may exchange data through shared memory. The tasks may be allocated to one or more processing units within a GPC **450** and instructions are scheduled for execution by at least one warp.

[0088] The work distribution unit **425** communicates with the one or more GPCs **450** via XBar **470**. The XBar **470** is an interconnect network that couples many of the units of the PPU **400** to other units of the PPU **400**. For example, the XBar **470** may be configured to couple the work distribution unit **425** to a particular GPC **450**. Although not shown explicitly, one or more other units of the PPU **400** may also be connected to the XBar **470** via the hub **430**.

[0089] The tasks are managed by the scheduler unit **420** and dispatched to a GPC **450** by the work distribution unit **425**. The GPC **450** is configured to process the task and generate results. The results may be consumed by other tasks within the GPC **450**, routed to a different GPC **450** via the XBar **470**, or stored in the memory **404**. The results can be written to the memory **404** via the memory partition units **480**, which implement a memory interface for reading and writing

data to/from the memory **404**. The results can be transmitted to another PPU **400** or CPU via the NVLink **410**. In an embodiment, the PPU **400** includes a number U of memory partition units **480** that is equal to the number of separate and distinct memory devices of the memory **404** coupled to the PPU **400**. Each GPC **450** may include a memory management unit to provide translation of virtual addresses into physical addresses, memory protection, and arbitration of memory requests. In an embodiment, the memory management unit provides one or more translation lookaside buffers (TLBs) for performing translation of virtual addresses into physical addresses in the memory **404**.

[0090] In an embodiment, the memory partition unit **480** includes a Raster Operations (ROP) unit, a level two (L2) cache, and a memory interface that is coupled to the memory **404**. The memory interface may implement 32, 64, 128, 1024-bit data buses, or the like, for high-speed data transfer. The PPU **400** may be connected to up to Y memory devices, such as high bandwidth memory stacks or graphics double-data-rate, version 5, synchronous dynamic random access memory, or other types of persistent storage. In an embodiment, the memory interface implements an HBM2 memory interface and Y equals half U . In an embodiment, the HBM2 memory stacks are located on the same physical package as the PPU **400**, providing substantial power and area savings compared with conventional GDDR5 SDRAM systems. In an embodiment, each HBM2 stack includes four memory dies and Y equals 4, with each HBM2 stack including two 128-bit channels per die for a total of 8 channels and a data bus width of 1024 bits.

[0091] In an embodiment, the memory **404** supports Single-Error Correcting Double-Error Detecting (SECCDED) Error Correction Code (ECC) to protect data. ECC provides higher reliability for compute applications that are sensitive to data corruption. Reliability is especially important in large-scale cluster computing environments where PPUs **400** process very large datasets and/or run applications for extended periods.

[0092] In an embodiment, the PPU **400** implements a multi-level memory hierarchy. In an embodiment, the memory partition unit **480** supports a unified memory to provide a single unified virtual address space for CPU and PPU **400** memory, enabling data sharing between virtual memory systems. In an embodiment the frequency of accesses by a PPU **400** to memory located on other processors is traced to ensure that memory pages are moved to the physical

memory of the PPU **400** that is accessing the pages more frequently. In an embodiment, the NVLink **410** supports address translation services allowing the PPU **400** to directly access a CPU's page tables and providing full access to CPU memory by the PPU **400**.

[0093] In an embodiment, copy engines transfer data between multiple PPUs **400** or between PPUs **400** and CPUs. The copy engines can generate page faults for addresses that are not mapped into the page tables. The memory partition unit **480** can then service the page faults, mapping the addresses into the page table, after which the copy engine can perform the transfer. In a conventional system, memory is pinned (e.g., non-pageable) for multiple copy engine operations between multiple processors, substantially reducing the available memory. With hardware page faulting, addresses can be passed to the copy engines without worrying if the memory pages are resident, and the copy process is transparent.

[0094] Data from the memory **404** or other system memory may be fetched by the memory partition unit **480** and stored in the L2 cache, which is located on-chip and is shared between the various GPCs **450**. As shown, each memory partition unit **480** includes a portion of the L2 cache associated with a corresponding memory **404**. Lower level caches may then be implemented in various units within the GPCs **450**. For example, each of the processing units within a GPC **450** may implement a level one (L1) cache. The L1 cache is private memory that is dedicated to a particular processing unit. The L2 cache is coupled to the memory interface **470** and the XBar **470** and data from the L2 cache may be fetched and stored in each of the L1 caches for processing.

[0095] In an embodiment, the processing units within each GPC **450** implement a SIMD (Single-Instruction, Multiple-Data) architecture where each thread in a group of threads (e.g., a warp) is configured to process a different set of data based on the same set of instructions. All threads in the group of threads execute the same instructions. In another embodiment, the processing unit implements a SIMT (Single-Instruction, Multiple Thread) architecture where each thread in a group of threads is configured to process a different set of data based on the same set of instructions, but where individual threads in the group of threads are allowed to diverge during execution. In an embodiment, a program counter, call stack, and execution state is maintained for each warp, enabling concurrency between warps and serial execution within warps when threads within the warp diverge. In another embodiment, a program counter, call

stack, and execution state is maintained for each individual thread, enabling equal concurrency between all threads, within and between warps. When execution state is maintained for each individual thread, threads executing the same instructions may be converged and executed in parallel for maximum efficiency.

[0096] Cooperative Groups is a programming model for organizing groups of communicating threads that allows developers to express the granularity at which threads are communicating, enabling the expression of richer, more efficient parallel decompositions. Cooperative launch APIs support synchronization amongst thread blocks for the execution of parallel algorithms. Conventional programming models provide a single, simple construct for synchronizing cooperating threads: a barrier across all threads of a thread block (e.g., the `syncthreads()` function). However, programmers would often like to define groups of threads at smaller than thread block granularities and synchronize within the defined groups to enable greater performance, design flexibility, and software reuse in the form of collective group-wide function interfaces.

[0097] Cooperative Groups enables programmers to define groups of threads explicitly at sub-block (e.g., as small as a single thread) and multi-block granularities, and to perform collective operations such as synchronization on the threads in a cooperative group. The programming model supports clean composition across software boundaries, so that libraries and utility functions can synchronize safely within their local context without having to make assumptions about convergence. Cooperative Groups primitives enable new patterns of cooperative parallelism, including producer-consumer parallelism, opportunistic parallelism, and global synchronization across an entire grid of thread blocks.

[0098] Each processing unit includes a large number (e.g., 128, etc.) of distinct processing cores (e.g., functional units) that may be fully-pipelined, single-precision, double-precision, and/or mixed precision and include a floating point arithmetic logic unit and an integer arithmetic logic unit. In an embodiment, the floating point arithmetic logic units implement the IEEE 754-2008 standard for floating point arithmetic. In an embodiment, the cores include 64 single-precision (32-bit) floating point cores, 64 integer cores, 32 double-precision (64-bit) floating point cores, and 8 tensor cores.

[0099] Tensor cores configured to perform matrix operations. In particular, the tensor cores are configured to perform deep learning matrix arithmetic, such as convolution operations for neural network training and inferencing. In an embodiment, each tensor core operates on a 4x4 matrix and performs a matrix multiply and accumulate operation $D=A \times B + C$, where A, B, C, and D are 4x4 matrices.

[00100] In an embodiment, the matrix multiply inputs A and B are 16-bit floating point matrices, while the accumulation matrices C and D may be 16-bit floating point or 32-bit floating point matrices. Tensor Cores operate on 16-bit floating point input data with 32-bit floating point accumulation. The 16-bit floating point multiply requires 64 operations and results in a full precision product that is then accumulated using 32-bit floating point addition with the other intermediate products for a 4x4x4 matrix multiply. In practice, Tensor Cores are used to perform much larger two-dimensional or higher dimensional matrix operations, built up from these smaller elements. An API, such as CUDA 9 C++ API, exposes specialized matrix load, matrix multiply and accumulate, and matrix store operations to efficiently use Tensor Cores from a CUDA-C++ program. At the CUDA level, the warp-level interface assumes 16x16 size matrices spanning all 32 threads of the warp.

[00101] Each processing unit may also comprise M special function units (SFUs) that perform special functions (e.g., attribute evaluation, reciprocal square root, and the like). In an embodiment, the SFUs may include a tree traversal unit configured to traverse a hierarchical tree data structure. In an embodiment, the SFUs may include texture unit configured to perform texture map filtering operations. In an embodiment, the texture units are configured to load texture maps (e.g., a 2D array of texels) from the memory **404** and sample the texture maps to produce sampled texture values for use in shader programs executed by the processing unit. In an embodiment, the texture maps are stored in shared memory that may comprise or include an L1 cache. The texture units implement texture operations such as filtering operations using mip-maps (e.g., texture maps of varying levels of detail). In an embodiment, each processing unit includes two texture units.

[00102] Each processing unit also comprises N load store units (LSUs) that implement load and store operations between the shared memory and the register file. Each processing unit includes an interconnect network that connects each of the cores to the register file and the LSU

to the register file, shared memory. In an embodiment, the interconnect network is a crossbar that can be configured to connect any of the cores to any of the registers in the register file and connect the LSUs to the register file and memory locations in shared memory.

[00103] The shared memory is an array of on-chip memory that allows for data storage and communication between the processing units and between threads within a processing unit. In an embodiment, the shared memory comprises 128KB of storage capacity and is in the path from each of the processing units to the memory partition unit **480**. The shared memory can be used to cache reads and writes. One or more of the shared memory, L1 cache, L2 cache, and memory **404** are backing stores.

[00104] Combining data cache and shared memory functionality into a single memory block provides the best overall performance for both types of memory accesses. The capacity is usable as a cache by programs that do not use shared memory. For example, if shared memory is configured to use half of the capacity, texture and load/store operations can use the remaining capacity. Integration within the shared memory enables the shared memory to function as a high-throughput conduit for streaming data while simultaneously providing high-bandwidth and low-latency access to frequently reused data.

[00105] When configured for general purpose parallel computation, a simpler configuration can be used compared with graphics processing. Specifically, fixed function graphics processing units, are bypassed, creating a much simpler programming model. In the general purpose parallel computation configuration, the work distribution unit **425** assigns and distributes blocks of threads directly to the processing units within the GPCs **450**. Threads execute the same program, using a unique thread ID in the calculation to ensure each thread generates unique results, using the processing unit(s) to execute the program and perform calculations, shared memory to communicate between threads, and the LSU to read and write global memory through the shared memory and the memory partition unit **480**. When configured for general purpose parallel computation, the processing units can also write commands that the scheduler unit **420** can use to launch new work on the processing units.

[00106] The PPUs **430** may each include, and/or be configured to perform functions of, one or more processing cores and/or components thereof, such as Tensor Cores (TCs), Tensor Processing Units (TPUs), Pixel Visual Cores (PVCs), Vision Processing Units (VPUs), Graphics

Processing Clusters (GPCs), Texture Processing Clusters (TPCs), Streaming Multiprocessors (SMs), Tree Traversal Units (TTUs), Artificial Intelligence Accelerators (AIAs), Deep Learning Accelerators (DLAs), Arithmetic-Logic Units (ALUs), Application-Specific Integrated Circuits (ASICs), Floating Point Units (FPUs), input/output (I/O) elements, peripheral component interconnect (PCI) or peripheral component interconnect express (PCIe) elements, and/or the like.

[00107] The PPU **400** may be included in a desktop computer, a laptop computer, a tablet computer, servers, supercomputers, a smart-phone (e.g., a wireless, hand-held device), personal digital assistant (PDA), a digital camera, a vehicle, a head mounted display, a hand-held electronic device, and the like. In an embodiment, the PPU **400** is embodied on a single semiconductor substrate. In another embodiment, the PPU **400** is included in a system-on-a-chip (SoC) along with one or more other devices such as additional PPUs **400**, the memory **404**, a reduced instruction set computer (RISC) CPU, a memory management unit (MMU), a digital-to-analog converter (DAC), and the like.

[00108] In an embodiment, the PPU **400** may be included on a graphics card that includes one or more memory devices. The graphics card may be configured to interface with a PCIe slot on a motherboard of a desktop computer. In yet another embodiment, the PPU **400** may be an integrated graphics processing unit (iGPU) or parallel processor included in the chipset of the motherboard.

Exemplary Computing System

[00109] Systems with multiple GPUs and CPUs are used in a variety of industries as developers expose and leverage more parallelism in applications such as artificial intelligence computing. High-performance GPU-accelerated systems with tens to many thousands of compute nodes are deployed in data centers, research facilities, and supercomputers to solve ever larger problems. As the number of processing devices within the high-performance systems increases, the communication and data transfer mechanisms need to scale to support the increased bandwidth.

[00110] Figure **5A** is a conceptual diagram of a processing system **500** implemented using the PPU **400** of Figure **3**, in accordance with an embodiment. The exemplary system **565** may be

configured to implement the method **130** shown in Figure **1B**, the method **148** shown in Figure **1D**, and/or the method **255** shown in Figure **2D**. The processing system **500** includes a CPU **530**, switch **510**, and multiple PPUs **400**, and respective memories **404**. The NVLink **410** provides high-speed communication links between each of the PPUs **400**. Although a particular number of NVLink **410** and interconnect **402** connections are illustrated in Figure **5A**, the number of connections to each PPU **400** and the CPU **530** may vary. The switch **510** interfaces between the interconnect **402** and the CPU **530**. The PPUs **400**, memories **404**, and NVLinks **410** may be situated on a single semiconductor platform to form a parallel processing module **525**. In an embodiment, the switch **510** supports two or more protocols to interface between various different connections and/or links.

[00111] In another embodiment (not shown), the NVLink **410** provides one or more high-speed communication links between each of the PPUs **400** and the CPU **530** and the switch **510** interfaces between the interconnect **402** and each of the PPUs **400**. The PPUs **400**, memories **404**, and interconnect **402** may be situated on a single semiconductor platform to form a parallel processing module **525**. In yet another embodiment (not shown), the interconnect **402** provides one or more communication links between each of the PPUs **400** and the CPU **530** and the switch **510** interfaces between each of the PPUs **400** using the NVLink **410** to provide one or more high-speed communication links between the PPUs **400**. In another embodiment (not shown), the NVLink **410** provides one or more high-speed communication links between the PPUs **400** and the CPU **530** through the switch **510**. In yet another embodiment (not shown), the interconnect **302** provides one or more communication links between each of the PPUs **400** directly. One or more of the NVLink **410** high-speed communication links may be implemented as a physical NVLink interconnect or either an on-chip or on-die interconnect using the same protocol as the NVLink **410**.

[00112] In the context of the present description, a single semiconductor platform may refer to a sole unitary semiconductor-based integrated circuit fabricated on a die or chip. It should be noted that the term single semiconductor platform may also refer to multi-chip modules with increased connectivity which simulate on-chip operation and make substantial improvements over utilizing a conventional bus implementation. Of course, the various circuits or devices may also be situated separately or in various combinations of semiconductor platforms per the

desires of the user. Alternately, the parallel processing module **525** may be implemented as a circuit board substrate and each of the PPUs **400** and/or memories **404** may be packaged devices. In an embodiment, the CPU **530**, switch **510**, and the parallel processing module **525** are situated on a single semiconductor platform.

[00113] In an embodiment, the signaling rate of each NVLink **410** is 20 to 25 Gigabits/second and each PPU **400** includes six NVLink **410** interfaces (as shown in Figure **5A**, five NVLink **410** interfaces are included for each PPU **400**). Each NVLink **410** provides a data transfer rate of 25 Gigabytes/second in each direction, with six links providing 300 Gigabytes/second. The NVLinks **410** can be used exclusively for PPU-to-PPU communication as shown in Figure **5A**, or some combination of PPU-to-PPU and PPU-to-CPU, when the CPU **530** also includes one or more NVLink **410** interfaces.

[00114] In an embodiment, the NVLink **410** allows direct load/store/atomic access from the CPU **530** to each PPU's **400** memory **404**. In an embodiment, the NVLink **410** supports coherency operations, allowing data read from the memories **404** to be stored in the cache hierarchy of the CPU **530**, reducing cache access latency for the CPU **530**. In an embodiment, the NVLink **410** includes support for Address Translation Services (ATS), allowing the PPU **400** to directly access page tables within the CPU **530**. One or more of the NVLinks **410** may also be configured to operate in a low-power mode.

[00115] Figure **5B** illustrates an exemplary system **565** in which the various architecture and/or functionality of the various previous embodiments may be implemented. The exemplary system **565** may be configured to implement the method **130** shown in Figure **1B**, the method **148** shown in Figure **1D**, and/or the method **255** shown in Figure **2D**.

[00116] As shown, a system **565** is provided including at least one central processing unit **530** that is connected to a communication bus **575**. The communication bus **575** may directly or indirectly couple one or more of the following devices: main memory **540**, network interface **535**, CPU(s) **530**, display device(s) **545**, input device(s) **560**, switch **510**, and parallel processing system **525**. The communication bus **575** may be implemented using any suitable protocol and may represent one or more links or busses, such as an address bus, a data bus, a control bus, or a combination thereof. The communication bus **575** may include one or more bus or link types, such as an industry standard architecture (ISA) bus, an extended industry standard architecture

(EISA) bus, a video electronics standards association (VESA) bus, a peripheral component interconnect (PCI) bus, a peripheral component interconnect express (PCIe) bus, HyperTransport, and/or another type of bus or link. In some embodiments, there are direct connections between components. As an example, the CPU(s) **530** may be directly connected to the main memory **540**. Further, the CPU(s) **530** may be directly connected to the parallel processing system **525**. Where there is direct, or point-to-point connection between components, the communication bus **575** may include a PCIe link to carry out the connection. In these examples, a PCI bus need not be included in the system **565**.

[00117] Although the various blocks of Figure **5C** are shown as connected via the communication bus **575** with lines, this is not intended to be limiting and is for clarity only. For example, in some embodiments, a presentation component, such as display device(s) **545**, may be considered an I/O component, such as input device(s) **560** (e.g., if the display is a touch screen). As another example, the CPU(s) **530** and/or parallel processing system **525** may include memory (e.g., the main memory **540** may be representative of a storage device in addition to the parallel processing system **525**, the CPUs **530**, and/or other components). In other words, the computing device of Figure **5C** is merely illustrative. Distinction is not made between such categories as “workstation,” “server,” “laptop,” “desktop,” “tablet,” “client device,” “mobile device,” “hand-held device,” “game console,” “electronic control unit (ECU),” “virtual reality system,” and/or other device or system types, as all are contemplated within the scope of the computing device of Figure **5C**.

[00118] The system **565** also includes a main memory **540**. Control logic (software) and data are stored in the main memory **540** which may take the form of a variety of computer-readable media. The computer-readable media may be any available media that may be accessed by the system **565**. The computer-readable media may include both volatile and nonvolatile media, and removable and non-removable media. By way of example, and not limitation, the computer-readable media may comprise computer-storage media and communication media.

[00119] The computer-storage media may include both volatile and nonvolatile media and/or removable and non-removable media implemented in any method or technology for storage of information such as computer-readable instructions, data structures, program modules, and/or

other data types. For example, the main memory **540** may store computer-readable instructions (e.g., that represent a program(s) and/or a program element(s), such as an operating system. Computer-storage media may include, but is not limited to, RAM, ROM, EEPROM, flash memory or other memory technology, CD-ROM, digital versatile disks (DVD) or other optical disk storage, magnetic cassettes, magnetic tape, magnetic disk storage or other magnetic storage devices, or any other medium which may be used to store the desired information and which may be accessed by system **565**. As used herein, computer storage media does not comprise signals per se.

[00120] The computer storage media may embody computer-readable instructions, data structures, program modules, and/or other data types in a modulated data signal such as a carrier wave or other transport mechanism and includes any information delivery media. The term “modulated data signal” may refer to a signal that has one or more of its characteristics set or changed in such a manner as to encode information in the signal. By way of example, and not limitation, the computer storage media may include wired media such as a wired network or direct-wired connection, and wireless media such as acoustic, RF, infrared and other wireless media. Combinations of any of the above should also be included within the scope of computer-readable media.

[00121] Computer programs, when executed, enable the system **565** to perform various functions. The CPU(s) **530** may be configured to execute at least some of the computer-readable instructions to control one or more components of the system **565** to perform one or more of the methods and/or processes described herein. The CPU(s) **530** may each include one or more cores (e.g., one, two, four, eight, twenty-eight, seventy-two, etc.) that are capable of handling a multitude of software threads simultaneously. The CPU(s) **530** may include any type of processor, and may include different types of processors depending on the type of system **565** implemented (e.g., processors with fewer cores for mobile devices and processors with more cores for servers). For example, depending on the type of system **565**, the processor may be an Advanced RISC Machines (ARM) processor implemented using Reduced Instruction Set Computing (RISC) or an x86 processor implemented using Complex Instruction Set Computing (CISC). The system **565** may include one or more CPUs **530** in addition to one or more microprocessors or supplementary co-processors, such as math co-processors.

[00122] In addition to or alternatively from the CPU(s) **530**, the parallel processing module **525** may be configured to execute at least some of the computer-readable instructions to control one or more components of the system **565** to perform one or more of the methods and/or processes described herein. The parallel processing module **525** may be used by the system **565** to render graphics (e.g., 3D graphics) or perform general purpose computations. For example, the parallel processing module **525** may be used for General-Purpose computing on GPUs (GPGPU). In embodiments, the CPU(s) **530** and/or the parallel processing module **525** may discretely or jointly perform any combination of the methods, processes and/or portions thereof.

[00123] The system **565** also includes input device(s) **560**, the parallel processing system **525**, and display device(s) **545**. The display device(s) **545** may include a display (e.g., a monitor, a touch screen, a television screen, a heads-up-display (HUD), other display types, or a combination thereof), speakers, and/or other presentation components. The display device(s) **545** may receive data from other components (e.g., the parallel processing system **525**, the CPU(s) **530**, etc.), and output the data (e.g., as an image, video, sound, etc.).

[00124] The network interface **535** may enable the system **565** to be logically coupled to other devices including the input devices **560**, the display device(s) **545**, and/or other components, some of which may be built in to (e.g., integrated in) the system **565**. Illustrative input devices **560** include a microphone, mouse, keyboard, joystick, game pad, game controller, satellite dish, scanner, printer, wireless device, etc. The input devices **560** may provide a natural user interface (NUI) that processes air gestures, voice, or other physiological inputs generated by a user. In some instances, inputs may be transmitted to an appropriate network element for further processing. An NUI may implement any combination of speech recognition, stylus recognition, facial recognition, biometric recognition, gesture recognition both on screen and adjacent to the screen, air gestures, head and eye tracking, and touch recognition (as described in more detail below) associated with a display of the system **565**. The system **565** may be include depth cameras, such as stereoscopic camera systems, infrared camera systems, RGB camera systems, touchscreen technology, and combinations of these, for gesture detection and recognition. Additionally, the system **565** may include accelerometers or gyroscopes (e.g., as part of an inertia measurement unit (IMU)) that enable detection of motion. In some examples,

the output of the accelerometers or gyroscopes may be used by the system **565** to render immersive augmented reality or virtual reality.

[00125] Further, the system **565** may be coupled to a network (e.g., a telecommunications network, local area network (LAN), wireless network, wide area network (WAN) such as the Internet, peer-to-peer network, cable network, or the like) through a network interface **535** for communication purposes. The system **565** may be included within a distributed network and/or cloud computing environment.

[00126] The network interface **535** may include one or more receivers, transmitters, and/or transceivers that enable the system **565** to communicate with other computing devices via an electronic communication network, included wired and/or wireless communications. The network interface **535** may include components and functionality to enable communication over any of a number of different networks, such as wireless networks (e.g., Wi-Fi, Z-Wave, Bluetooth, Bluetooth LE, ZigBee, etc.), wired networks (e.g., communicating over Ethernet or InfiniBand), low-power wide-area networks (e.g., LoRaWAN, SigFox, etc.), and/or the Internet.

[00127] The system **565** may also include a secondary storage (not shown). The secondary storage **610** includes, for example, a hard disk drive and/or a removable storage drive, representing a floppy disk drive, a magnetic tape drive, a compact disk drive, digital versatile disk (DVD) drive, recording device, universal serial bus (USB) flash memory. The removable storage drive reads from and/or writes to a removable storage unit in a well-known manner. The system **565** may also include a hard-wired power supply, a battery power supply, or a combination thereof (not shown). The power supply may provide power to the system **565** to enable the components of the system **565** to operate.

[00128] Each of the foregoing modules and/or devices may even be situated on a single semiconductor platform to form the system **565**. Alternately, the various modules may also be situated separately or in various combinations of semiconductor platforms per the desires of the user.

[00129] While various embodiments have been described above, it should be understood that they have been presented by way of example only, and not limitation. Thus, the breadth and scope of a preferred embodiment should not be limited by any of the above-described exemplary

embodiments, but should be defined only in accordance with the following claims and their equivalents.

Machine Learning

[00130] Deep neural networks (DNNs) developed on processors, such as the PPU **400** have been used for diverse use cases, from self-driving cars to faster drug development, from automatic image captioning in online image databases to smart real-time language translation in video chat applications. Deep learning is a technique that models the neural learning process of the human brain, continually learning, continually getting smarter, and delivering more accurate results more quickly over time. A child is initially taught by an adult to correctly identify and classify various shapes, eventually being able to identify shapes without any coaching. Similarly, a deep learning or neural learning system needs to be trained in object recognition and classification for it get smarter and more efficient at identifying basic objects, occluded objects, etc., while also assigning context to objects.

[00131] At the simplest level, neurons in the human brain look at various inputs that are received, importance levels are assigned to each of these inputs, and output is passed on to other neurons to act upon. An artificial neuron or perceptron is the most basic model of a neural network. In one example, a perceptron may receive one or more inputs that represent various features of an object that the perceptron is being trained to recognize and classify, and each of these features is assigned a certain weight based on the importance of that feature in defining the shape of an object.

[00132] A deep neural network (DNN) model includes multiple layers of many connected nodes (e.g., perceptrons, Boltzmann machines, radial basis functions, convolutional layers, etc.) that can be trained with enormous amounts of input data to quickly solve complex problems with high accuracy. In one example, a first layer of the DNN model breaks down an input image of an automobile into various sections and looks for basic patterns such as lines and angles. The second layer assembles the lines to look for higher level patterns such as wheels, windshields, and mirrors. The next layer identifies the type of vehicle, and the final few layers generate a label for the input image, identifying the model of a specific automobile brand.

[00133] Once the DNN is trained, the DNN can be deployed and used to identify and classify objects or patterns in a process known as inference. Examples of inference (the process through which a DNN extracts useful information from a given input) include identifying handwritten numbers on checks deposited into ATM machines, identifying images of friends in photos, delivering movie recommendations to over fifty million users, identifying and classifying different types of automobiles, pedestrians, and road hazards in driverless cars, or translating human speech in real-time.

[00134] During training, data flows through the DNN in a forward propagation phase until a prediction is produced that indicates a label corresponding to the input. If the neural network does not correctly label the input, then errors between the correct label and the predicted label are analyzed, and the weights are adjusted for each feature during a backward propagation phase until the DNN correctly labels the input and other inputs in a training dataset. Training complex neural networks requires massive amounts of parallel computing performance, including floating-point multiplications and additions that are supported by the PPU **400**. Inferencing is less compute-intensive than training, being a latency-sensitive process where a trained neural network is applied to new inputs it has not seen before to classify images, translate speech, and generally infer new information.

[00135] Neural networks rely heavily on matrix math operations, and complex multi-layered networks require tremendous amounts of floating-point performance and bandwidth for both efficiency and speed. With thousands of processing cores, optimized for matrix math operations, and delivering tens to hundreds of TFLOPS of performance, the PPU **400** is a computing platform capable of delivering performance required for deep neural network-based artificial intelligence and machine learning applications.

[00136] Furthermore, images generated applying one or more of the techniques disclosed herein may be used to train, test, or certify DNNs used to recognize objects and environments in the real world. Such images may include scenes of roadways, factories, buildings, urban settings, rural settings, humans, animals, and any other physical object or real-world setting. Such images may be used to train, test, or certify DNNs that are employed in machines or robots to manipulate, handle, or modify physical objects in the real world. Furthermore, such images may be used to train, test, or certify DNNs that are employed in autonomous vehicles to navigate

and move the vehicles through the real world. Additionally, images generated applying one or more of the techniques disclosed herein may be used to convey information to users of such machines, robots, and vehicles.

Graphics Processing Pipeline

[00137] In an embodiment, the PPU **400** comprises a graphics processing unit (GPU). The PPU **400** is configured to receive commands that specify shader programs for processing graphics data. Graphics data may be defined as a set of primitives such as points, lines, triangles, quads, triangle strips, and the like. Typically, a primitive includes data that specifies a number of vertices for the primitive (e.g., in a model-space coordinate system) as well as attributes associated with each vertex of the primitive. The PPU **400** can be configured to process the graphics primitives to generate a frame buffer (e.g., pixel data for each of the pixels of the display).

[00138] An application writes model data for a scene (e.g., a collection of vertices and attributes) to a memory such as a system memory or memory **404**. The model data defines each of the objects that may be visible on a display. The application then makes an API call to the driver kernel that requests the model data to be rendered and displayed. The driver kernel reads the model data and writes commands to the one or more streams to perform operations to process the model data. The commands may reference different shader programs to be implemented on the processing units of the PPU **400** including one or more of a vertex shader, hull shader, domain shader, geometry shader, and a pixel shader. For example, one or more of the processing units may be configured to execute a vertex shader program that processes a number of vertices defined by the model data. In an embodiment, the different processing units may be configured to execute different shader programs concurrently. For example, a first subset of processing units may be configured to execute a vertex shader program while a second subset of processing units may be configured to execute a pixel shader program. The first subset of processing units processes vertex data to produce processed vertex data and writes the processed vertex data to the L2 cache and/or the memory **404**. After the processed vertex data is rasterized (e.g., transformed from three-dimensional data into two-dimensional data in screen space) to produce fragment data, the second subset of processing units executes a pixel shader to produce processed fragment data, which is then blended with other processed fragment data and

written to the frame buffer in memory **404**. The vertex shader program and pixel shader program may execute concurrently, processing different data from the same scene in a pipelined fashion until all of the model data for the scene has been rendered to the frame buffer. Then, the contents of the frame buffer are transmitted to a display controller for display on a display device.

[00139] Figure **6A** is a conceptual diagram of a graphics processing pipeline **600** implemented by the PPU **400** of Figure **4**, in accordance with an embodiment. The graphics processing pipeline **600** is an abstract flow diagram of the processing steps implemented to generate 2D computer-generated images from 3D geometry data. As is well-known, pipeline architectures may perform long latency operations more efficiently by splitting up the operation into a plurality of stages, where the output of each stage is coupled to the input of the next successive stage. Thus, the graphics processing pipeline **600** receives input data **601** that is transmitted from one stage to the next stage of the graphics processing pipeline **600** to generate output data **602**. In an embodiment, the graphics processing pipeline **600** may represent a graphics processing pipeline defined by the OpenGL[®] API. As an option, the graphics processing pipeline **600** may be implemented in the context of the functionality and architecture of the previous Figures and/or any subsequent Figure(s).

[00140] As shown in Figure **6A**, the graphics processing pipeline **600** comprises a pipeline architecture that includes a number of stages. The stages include, but are not limited to, a data assembly stage **610**, a vertex shading stage **620**, a primitive assembly stage **630**, a geometry shading stage **640**, a viewport scale, cull, and clip (VSCC) stage **650**, a rasterization stage **660**, a fragment shading stage **670**, and a raster operations stage **680**. In an embodiment, the input data **601** comprises commands that configure the processing units to implement the stages of the graphics processing pipeline **600** and geometric primitives (e.g., points, lines, triangles, quads, triangle strips or fans, etc.) to be processed by the stages. The output data **602** may comprise pixel data (e.g., color data) that is copied into a frame buffer or other type of surface data structure in a memory.

[00141] The data assembly stage **610** receives the input data **601** that specifies vertex data for high-order surfaces, primitives, or the like. The data assembly stage **610** collects the vertex data in a temporary storage or queue, such as by receiving a command from the host processor

that includes a pointer to a buffer in memory and reading the vertex data from the buffer. The vertex data is then transmitted to the vertex shading stage **620** for processing.

[00142] The vertex shading stage **620** processes vertex data by performing a set of operations (e.g., a vertex shader or a program) once for each of the vertices. Vertices may be, e.g., specified as a 4-coordinate vector (e.g., $\langle x, y, z, w \rangle$) associated with one or more vertex attributes (e.g., color, texture coordinates, surface normal, etc.). The vertex shading stage **620** may manipulate individual vertex attributes such as position, color, texture coordinates, and the like. In other words, the vertex shading stage **620** performs operations on the vertex coordinates or other vertex attributes associated with a vertex. Such operations commonly including lighting operations (e.g., modifying color attributes for a vertex) and transformation operations (e.g., modifying the coordinate space for a vertex). For example, vertices may be specified using coordinates in an object-coordinate space, which are transformed by multiplying the coordinates by a matrix that translates the coordinates from the object-coordinate space into a world space or a normalized-device-coordinate (NCD) space. The vertex shading stage **620** generates transformed vertex data that is transmitted to the primitive assembly stage **630**.

[00143] The primitive assembly stage **630** collects vertices output by the vertex shading stage **620** and groups the vertices into geometric primitives for processing by the geometry shading stage **640**. For example, the primitive assembly stage **630** may be configured to group every three consecutive vertices as a geometric primitive (e.g., a triangle) for transmission to the geometry shading stage **640**. In some embodiments, specific vertices may be reused for consecutive geometric primitives (e.g., two consecutive triangles in a triangle strip may share two vertices). The primitive assembly stage **630** transmits geometric primitives (e.g., a collection of associated vertices) to the geometry shading stage **640**.

[00144] The geometry shading stage **640** processes geometric primitives by performing a set of operations (e.g., a geometry shader or program) on the geometric primitives. Tessellation operations may generate one or more geometric primitives from each geometric primitive. In other words, the geometry shading stage **640** may subdivide each geometric primitive into a finer mesh of two or more geometric primitives for processing by the rest of the graphics processing pipeline **600**. The geometry shading stage **640** transmits geometric primitives to the viewport SCC stage **650**.

[00145] In an embodiment, the graphics processing pipeline **600** may operate within a streaming multiprocessor and the vertex shading stage **620**, the primitive assembly stage **630**, the geometry shading stage **640**, the fragment shading stage **670**, and/or hardware/software associated therewith, may sequentially perform processing operations. Once the sequential processing operations are complete, in an embodiment, the viewport SCC stage **650** may utilize the data. In an embodiment, primitive data processed by one or more of the stages in the graphics processing pipeline **600** may be written to a cache (e.g. L1 cache, a vertex cache, etc.). In this case, in an embodiment, the viewport SCC stage **650** may access the data in the cache. In an embodiment, the viewport SCC stage **650** and the rasterization stage **660** are implemented as fixed function circuitry.

[00146] The viewport SCC stage **650** performs viewport scaling, culling, and clipping of the geometric primitives. Each surface being rendered to is associated with an abstract camera position. The camera position represents a location of a viewer looking at the scene and defines a viewing frustum that encloses the objects of the scene. The viewing frustum may include a viewing plane, a rear plane, and four clipping planes. Any geometric primitive entirely outside of the viewing frustum may be culled (e.g., discarded) because the geometric primitive will not contribute to the final rendered scene. Any geometric primitive that is partially inside the viewing frustum and partially outside the viewing frustum may be clipped (e.g., transformed into a new geometric primitive that is enclosed within the viewing frustum). Furthermore, geometric primitives may each be scaled based on a depth of the viewing frustum. All potentially visible geometric primitives are then transmitted to the rasterization stage **660**.

[00147] The rasterization stage **660** converts the 3D geometric primitives into 2D fragments (e.g. capable of being utilized for display, etc.). The rasterization stage **660** may be configured to utilize the vertices of the geometric primitives to setup a set of plane equations from which various attributes can be interpolated. The rasterization stage **660** may also compute a coverage mask for a plurality of pixels that indicates whether one or more sample locations for the pixel intercept the geometric primitive. In an embodiment, z-testing may also be performed to determine if the geometric primitive is occluded by other geometric primitives that have already been rasterized. The rasterization stage **660** generates fragment data (e.g., interpolated vertex

attributes associated with a particular sample location for each covered pixel) that are transmitted to the fragment shading stage **670**.

[00148] The fragment shading stage **670** processes fragment data by performing a set of operations (e.g., a fragment shader or a program) on each of the fragments. The fragment shading stage **670** may generate pixel data (e.g., color values) for the fragment such as by performing lighting operations or sampling texture maps using interpolated texture coordinates for the fragment. The fragment shading stage **670** generates pixel data that is transmitted to the raster operations stage **680**.

[00149] The raster operations stage **680** may perform various operations on the pixel data such as performing alpha tests, stencil tests, and blending the pixel data with other pixel data corresponding to other fragments associated with the pixel. When the raster operations stage **680** has finished processing the pixel data (e.g., the output data **602**), the pixel data may be written to a render target such as a frame buffer, a color buffer, or the like.

[00150] It will be appreciated that one or more additional stages may be included in the graphics processing pipeline **600** in addition to or in lieu of one or more of the stages described above. Various implementations of the abstract graphics processing pipeline may implement different stages. Furthermore, one or more of the stages described above may be excluded from the graphics processing pipeline in some embodiments (such as the geometry shading stage **640**). Other types of graphics processing pipelines are contemplated as being within the scope of the present disclosure. Furthermore, any of the stages of the graphics processing pipeline **600** may be implemented by one or more dedicated hardware units within a graphics processor such as PPU **400**. Other stages of the graphics processing pipeline **600** may be implemented by programmable hardware units such as the processing unit within the PPU **400**.

[00151] The graphics processing pipeline **600** may be implemented via an application executed by a host processor, such as a CPU. In an embodiment, a device driver may implement an application programming interface (API) that defines various functions that can be utilized by an application in order to generate graphical data for display. The device driver is a software program that includes a plurality of instructions that control the operation of the PPU **400**. The API provides an abstraction for a programmer that lets a programmer utilize

specialized graphics hardware, such as the PPU 400, to generate the graphical data without requiring the programmer to utilize the specific instruction set for the PPU 400. The application may include an API call that is routed to the device driver for the PPU 400. The device driver interprets the API call and performs various operations to respond to the API call. In some instances, the device driver may perform operations by executing instructions on the CPU. In other instances, the device driver may perform operations, at least in part, by launching operations on the PPU 400 utilizing an input/output interface between the CPU and the PPU 400. In an embodiment, the device driver is configured to implement the graphics processing pipeline 600 utilizing the hardware of the PPU 400.

[00152] Various programs may be executed within the PPU 400 in order to implement the various stages of the graphics processing pipeline 600. For example, the device driver may launch a kernel on the PPU 400 to perform the vertex shading stage 620 on one processing unit (or multiple processing units). The device driver (or the initial kernel executed by the PPU 400) may also launch other kernels on the PPU 400 to perform other stages of the graphics processing pipeline 600, such as the geometry shading stage 640 and the fragment shading stage 670. In addition, some of the stages of the graphics processing pipeline 600 may be implemented on fixed unit hardware such as a rasterizer or a data assembler implemented within the PPU 400. It will be appreciated that results from one kernel may be processed by one or more intervening fixed function hardware units before being processed by a subsequent kernel on a processing unit.

[00153] Images generated applying one or more of the techniques disclosed herein may be displayed on a monitor or other display device. In some embodiments, the display device may be coupled directly to the system or processor generating or rendering the images. In other embodiments, the display device may be coupled indirectly to the system or processor such as via a network. Examples of such networks include the Internet, mobile telecommunications networks, a WIFI network, as well as any other wired and/or wireless networking system. When the display device is indirectly coupled, the images generated by the system or processor may be streamed over the network to the display device. Such streaming allows, for example, video games or other applications, which render images, to be executed on a server, in a data center, or in a cloud-based computing environment and the rendered images to be transmitted

and displayed on one or more user devices (such as a computer, video game console, smartphone, other mobile device, *etc.*) that are physically separate from the server or data center. Hence, the techniques disclosed herein can be applied to enhance the images that are streamed and to enhance services that stream images such as NVIDIA GeForce Now (GFN), Google Stadia, and the like.

Example Game Streaming System

[00154] Figure **6B** is an example system diagram for a game streaming system **605**, in accordance with some embodiments of the present disclosure. Figure **6B** includes game server(s) **603** (which may include similar components, features, and/or functionality to the example processing system **500** of Figure **5A** and/or exemplary system **565** of Figure **5B**), client device(s) **604** (which may include similar components, features, and/or functionality to the example processing system **500** of Figure **5A** and/or exemplary system **565** of Figure **5B**), and network(s) **606** (which may be similar to the network(s) described herein). In some embodiments of the present disclosure, the system **605** may be implemented.

[00155] In the system **605**, for a game session, the client device(s) **604** may only receive input data in response to inputs to the input device(s) **626**, transmit the input data to the game server(s) **603**, receive encoded display data from the game server(s) **603**, and display the display data on the display **624**. As such, the more computationally intense computing and processing is offloaded to the game server(s) **603** (e.g., rendering – in particular ray or path tracing – for graphical output of the game session is executed by the GPU(s) **615** of the game server(s) **603**). In other words, the game session is streamed to the client device(s) **604** from the game server(s) **603**, thereby reducing the requirements of the client device(s) **604** for graphics processing and rendering.

[00156] For example, with respect to an instantiation of a game session, a client device **604** may be displaying a frame of the game session on the display **624** based on receiving the display data from the game server(s) **603**. The client device **604** may receive an input to one of the input device(s) **626** and generate input data in response. The client device **604** may transmit the input data to the game server(s) **603** via the communication interface **621** and over the network(s) **606** (e.g., the Internet), and the game server(s) **603** may receive the input data via the communication interface **618**. The CPU(s) **608** may receive the input data, process the input

data, and transmit data to the GPU(s) **615** that causes the GPU(s) **615** to generate a rendering of the game session. For example, the input data may be representative of a movement of a character of the user in a game, firing a weapon, reloading, passing a ball, turning a vehicle, etc. The rendering component **612** may render the game session (e.g., representative of the result of the input data) and the render capture component **614** may capture the rendering of the game session as display data (e.g., as image data capturing the rendered frame of the game session). The rendering of the game session may include ray or path-traced lighting and/or shadow effects, computed using one or more parallel processing units – such as GPUs **615**, which may further employ the use of one or more dedicated hardware accelerators or processing cores to perform ray or path-tracing techniques – of the game server(s) **603**. The encoder **616** may then encode the display data to generate encoded display data and the encoded display data may be transmitted to the client device **604** over the network(s) **606** via the communication interface **618**. The client device **604** may receive the encoded display data via the communication interface **621** and the decoder **622** may decode the encoded display data to generate the display data. The client device **604** may then display the display data via the display **624**.

Example Network Environments

[00157] Network environments suitable for use in implementing embodiments of the disclosure may include one or more client devices, servers, network attached storage (NAS), other backend devices, and/or other device types. The client devices, servers, and/or other device types (e.g., each device) may be implemented on one or more instances of the processing system **500** of Figure **5A** and/or exemplary system **565** of Figure **5B** – e.g., each device may include similar components, features, and/or functionality of the processing system **500** and/or exemplary system **565**.

[00158] Components of a network environment may communicate with each other via a network(s), which may be wired, wireless, or both. The network may include multiple networks, or a network of networks. By way of example, the network may include one or more Wide Area Networks (WANs), one or more Local Area Networks (LANs), one or more public networks such as the Internet and/or a public switched telephone network (PSTN), and/or one or more private networks. Where the network includes a wireless telecommunications network,

components such as a base station, a communications tower, or even access points (as well as other components) may provide wireless connectivity.

[00159] Compatible network environments may include one or more peer-to-peer network environments – in which case a server may not be included in a network environment – and one or more client-server network environments – in which case one or more servers may be included in a network environment. In peer-to-peer network environments, functionality described herein with respect to a server(s) may be implemented on any number of client devices.

[00160] In at least one embodiment, a network environment may include one or more cloud-based network environments, a distributed computing environment, a combination thereof, etc. A cloud-based network environment may include a framework layer, a job scheduler, a resource manager, and a distributed file system implemented on one or more of servers, which may include one or more core network servers and/or edge servers. A framework layer may include a framework to support software of a software layer and/or one or more application(s) of an application layer. The software or application(s) may respectively include web-based service software or applications. In embodiments, one or more of the client devices may use the web-based service software or applications (e.g., by accessing the service software and/or applications via one or more application programming interfaces (APIs)). The framework layer may be, but is not limited to, a type of free and open-source software web application framework such as that may use a distributed file system for large-scale data processing (e.g., "big data").

[00161] A cloud-based network environment may provide cloud computing and/or cloud storage that carries out any combination of computing and/or data storage functions described herein (or one or more portions thereof). Any of these various functions may be distributed over multiple locations from central or core servers (e.g., of one or more data centers that may be distributed across a state, a region, a country, the globe, etc.). If a connection to a user (e.g., a client device) is relatively close to an edge server(s), a core server(s) may designate at least a portion of the functionality to the edge server(s). A cloud-based network environment may be private (e.g., limited to a single organization), may be public (e.g., available to many organizations), and/or a combination thereof (e.g., a hybrid cloud environment).

[00162] The client device(s) may include at least some of the components, features, and functionality of the example processing system **500** of Figure **5B** and/or exemplary system **565** of Figure **5C**. By way of example and not limitation, a client device may be embodied as a Personal Computer (PC), a laptop computer, a mobile device, a smartphone, a tablet computer, a smart watch, a wearable computer, a Personal Digital Assistant (PDA), an MP3 player, a virtual reality headset, a Global Positioning System (GPS) or device, a video player, a video camera, a surveillance device or system, a vehicle, a boat, a flying vessel, a virtual machine, a drone, a robot, a handheld communications device, a hospital device, a gaming device or system, an entertainment system, a vehicle computer system, an embedded system controller, a remote control, an appliance, a consumer electronic device, a workstation, an edge device, any combination of these delineated devices, or any other suitable device.

[00163] It is noted that the techniques described herein may be embodied in executable instructions stored in a computer readable medium for use by or in connection with a processor-based instruction execution machine, system, apparatus, or device. It will be appreciated by those skilled in the art that, for some embodiments, various types of computer-readable media can be included for storing data. As used herein, a "computer-readable medium" includes one or more of any suitable media for storing the executable instructions of a computer program such that the instruction execution machine, system, apparatus, or device may read (or fetch) the instructions from the computer-readable medium and execute the instructions for carrying out the described embodiments. Suitable storage formats include one or more of an electronic, magnetic, optical, and electromagnetic format. A non-exhaustive list of conventional exemplary computer-readable medium includes: a portable computer diskette; a random-access memory (RAM); a read-only memory (ROM); an erasable programmable read only memory (EPROM); a flash memory device; and optical storage devices, including a portable compact disc (CD), a portable digital video disc (DVD), and the like.

[00164] It should be understood that the arrangement of components illustrated in the attached Figures are for illustrative purposes and that other arrangements are possible. For example, one or more of the elements described herein may be realized, in whole or in part, as an electronic hardware component. Other elements may be implemented in software, hardware, or a combination of software and hardware. Moreover, some or all of these other elements may be

combined, some may be omitted altogether, and additional components may be added while still achieving the functionality described herein. Thus, the subject matter described herein may be embodied in many different variations, and all such variations are contemplated to be within the scope of the claims.

[00165] To facilitate an understanding of the subject matter described herein, many aspects are described in terms of sequences of actions. It will be recognized by those skilled in the art that the various actions may be performed by specialized circuits or circuitry, by program instructions being executed by one or more processors, or by a combination of both. The description herein of any sequence of actions is not intended to imply that the specific order described for performing that sequence must be followed. All methods described herein may be performed in any suitable order unless otherwise indicated herein or otherwise clearly contradicted by context.

[00166] The use of the terms "a" and "an" and "the" and similar references in the context of describing the subject matter (particularly in the context of the following claims) are to be construed to cover both the singular and the plural, unless otherwise indicated herein or clearly contradicted by context. The use of the term "at least one" followed by a list of one or more items (for example, "at least one of A and B") is to be construed to mean one item selected from the listed items (A or B) or any combination of two or more of the listed items (A and B), unless otherwise indicated herein or clearly contradicted by context. Furthermore, the foregoing description is for the purpose of illustration only, and not for the purpose of limitation, as the scope of protection sought is defined by the claims as set forth hereinafter together with any equivalents thereof. The use of any and all examples, or exemplary language (e.g., "such as") provided herein, is intended merely to better illustrate the subject matter and does not pose a limitation on the scope of the subject matter unless otherwise claimed. The use of the term "based on" and other like phrases indicating a condition for bringing about a result, both in the claims and in the written description, is not intended to foreclose any other conditions that bring about that result. No language in the specification should be construed as indicating any non-claimed element as essential to the practice of the invention as claimed.